

week1

Rajesh Kalakoti

2023-08-03

- Packages
 - devtools
 - tidyverse
 - [here]

```
# Include the script from the R directory
project_path <- here()
source(here("R", "utils.R"))
source(here("R", "distance_functions.R"))
```

Clustering

Given a clustering $C = \{C_1, C_2, \dots, C_k\}$, we need some scoring function that evaluates its quality or goodness. This sum of squared errors scoring function is defined as:

$$W(C) = \frac{1}{2} \sum_{k=1}^K \sum_{i: C(i)=k} \|x_i - \bar{x}_k\|^2$$

The goal is to find the clustering that minimizes:

$$C^* = \arg \min_C \{W(c)\}$$

K-means employs a greedy iterative approach to find a clustering that minimizes loss function.

Note that the `echo = FALSE` parameter was added to the code chunk to prevent printing of the R code that generated the plot.

```
euclidean_dist <- function(point1, point2) {
  squared_diff <- (point1 - point2)^2
  sqrt(sum(squared_diff))
}

x <- y <- seq(-1, 1, length = 20)
grid <- expand.grid(x = x, y = y) # Create a grid of points
z <- matrix(0, nrow = length(x), ncol = length(y)) # Initialize the z matrix

for (i in 1:length(x)) {
  for (j in 1:length(y)) {
    z[i, j] <- euclidean_dist(c(x[i], y[j]), c(0, 0))
  }
}

persp(x, y, z,
```

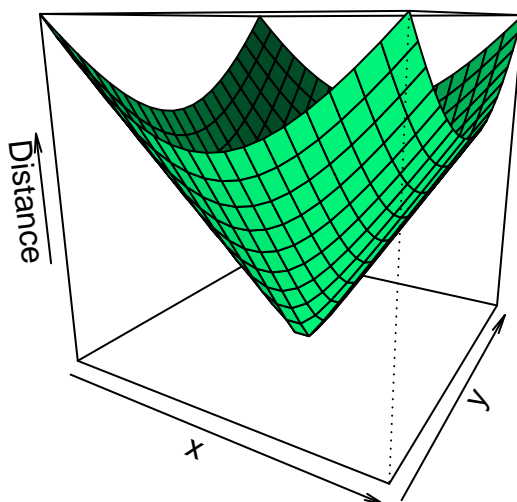
Algorithm 1: K-means Algorithm

Data: D, k, ε

```
1 K-means( $D, k, \varepsilon$ ):  
2  $t \leftarrow 0$ ;  
3 Randomly initialize  $k$  centroids:  $\mu_1^t, \mu_2^t, \dots, \mu_n^t \in \mathbb{R}^d$ ;  
4 repeat  
5    $t \leftarrow t + 1$ ;  
6    $C_i \leftarrow \emptyset$  for all  $i = 1, \dots, k$   
7   /* Cluster assignment step */  
8   for  $x_j \in D$  do  
9      $i^* \leftarrow \operatorname{argmin}_i \{\|x_j - \mu_i^{t-1}\|^2\}$ ;  
10    /* assign  $x_j$  to closest centroid */  
11     $C_{i^*} \leftarrow C_{i^*} \cup \{x_j\}$ ;  
12  end  
13  for  $i = 1, \dots, k$  do  
14     $\mu_i^t \leftarrow \frac{1}{|C_i|} \sum_{x_j \in C_i} X_j$   
15  end  
16 until  $\sum_{i=1}^k \|\mu_i^t - \mu_i^{t-1}\|^2 \leq \varepsilon$ ;
```

```
main = "3D Plot of Euclidean Distance",  
zlab = "Distance",  
theta = 30, phi = 15,  
col = "springgreen", shade = 0.5)
```

3D Plot of Euclidean Distance



```

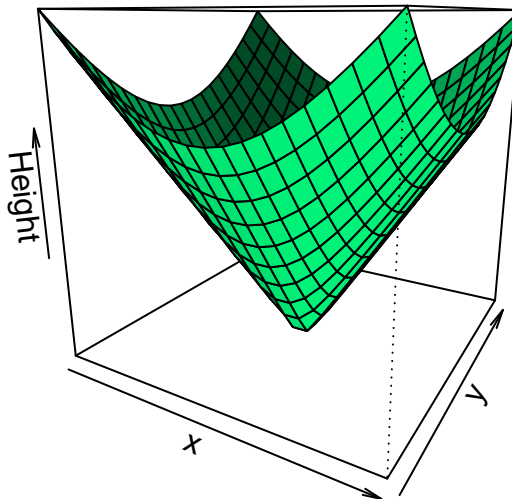
cone1 <- function(x, y){
  sqrt(x^2+y^2)
}

x <- y <- seq(-1, 1, length= 20)
z <- outer(x, y, cone1)

persp(x, y, z,
  main="Perspective Plot of a Cone",
  zlab = "Height",
  theta = 30, phi = 15,
  col = "springgreen", shade = 0.5)

```

Perspective Plot of a Cone



```

euclidean_dist <- function(point1, point2) {
  squared_diff <- (point1 - point2)^2
  sqrt(sum(squared_diff))
}

manhattan_distance <- function(point1, point2) {
  if (length(point1) != length(point2)) {
    stop("Both points should have the same number of dimensions.")
  }

  abs_diff <- abs(point1 - point2)
  distance <- sum(abs_diff)
  return(distance)
}

```

```

}

x <- y <- seq(-1, 1, length = 20)
grid <- expand.grid(x = x, y = y) # Create a grid of points
z_euclidean <- matrix(0, nrow = length(x), ncol = length(y)) # Initialize the z matrix for Euclidean d
z_manhattan <- matrix(0, nrow = length(x), ncol = length(y)) # Initialize the z matrix for Manhattan

for (i in 1:length(x)) {
  for (j in 1:length(y)) {
    z_euclidean[i, j] <- euclidean_dist(c(x[i], y[j]), c(0, 0))
    z_manhattan[i, j] <- manhattan_distance(c(x[i], y[j]), c(0, 0))
  }
}

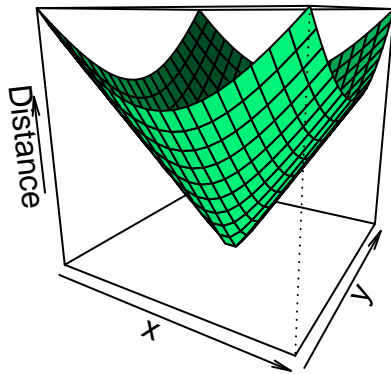
# Create a layout of subplots to show both Euclidean and Manhattan distances
par(mfrow = c(1, 2))

# Plot for Euclidean distance
persp(x, y, z_euclidean,
      main = "3D Plot of Euclidean Distance",
      zlab = "Distance",
      theta = 30, phi = 15,
      col = "springgreen", shade = 0.5)

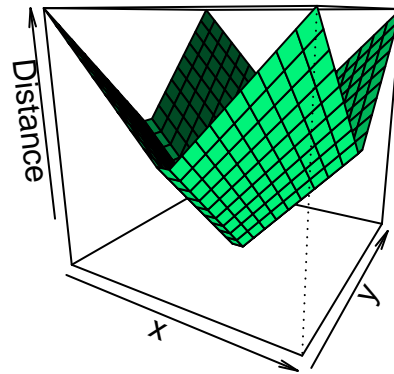
# Plot for Manhattan distance
persp(x, y, z_manhattan,
      main = "3D Plot of Manhattan Distance",
      zlab = "Distance",
      theta = 30, phi = 15,
      col = "springgreen", shade = 0.5)

```

3D Plot of Euclidean Distance



3D Plot of Manhattan Distance



```
# Reset the layout
par(mfrow = c(1, 1))
```

```
euclidean_dist <- function(point1, point2) {
  squared_diff <- (point1 - point2)^2
  sqrt(sum(squared_diff))
}
```

```
manhattan_distance <- function(point1, point2) {
  if (length(point1) != length(point2)) {
    stop("Both points should have the same number of dimensions.")
  }
}
```

```
abs_diff <- abs(point1 - point2)
distance <- sum(abs_diff)
return(distance)
}
```

```
x <- y <- seq(-5, 5, length = 20)
grid <- expand.grid(x = x, y = y) # Create a grid of points
```

```
z_euclidean <- matrix(0, nrow = length(x), ncol = length(y)) # Initialize the z matrix for Euclidean d
z_manhattan <- matrix(0, nrow = length(x), ncol = length(y)) # Initialize the z matrix for Manhattan
```

```
for (i in 1:length(x)) {
  for (j in 1:length(y)) {
```

```

    z_euclidean[i, j] <- euclidean_dist(c(x[i], y[j]), c(0, 0))
    z_manhattan[i, j] <- manhattan_distance(c(x[i], y[j]), c(0, 0))
  }
}

# Combine the distances and choose different colors for each
combined_distances <- z_euclidean + z_manhattan
color_palette <- colorRampPalette(c("blue", "green"))(100) # Choose colors for mapping distances

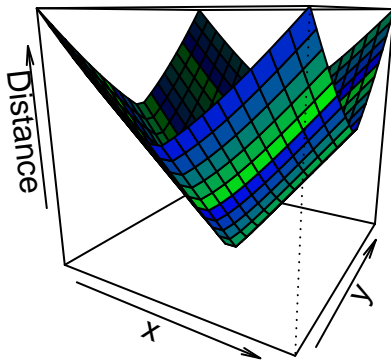
# Create a layout of subplots
layout(matrix(c(1, 2), nrow = 1))

# Plot both distances on the same 3D plane with different colors
persp(x, y, combined_distances,
      main = "3D Plot of Combined Distances",
      zlab = "Distance",
      theta = 30, phi = 15,
      col = color_palette, shade = 0.5)

# Reset the layout
layout(1)

```

3D Plot of Combined Distances



```

# Include the script from the R directory
project_path <- here()
source(here("R", "utils.R"))

```

```
?entropy
```

```
## No documentation for 'entropy' in specified packages and libraries:  
## you could try '??entropy'
```